

NPTEL Research Methodology — Week 6 Research-Grade Notes

I’ve treated **Week 6 / Module 6, Lectures 31–36** as the backbone and applied the master prompt you provided: preserve the lecturer's framework, but distinguish lecture claims from methodological precision, verification, and additional knowledge. Your master prompt explicitly requires that the transcript be treated as a source of claims rather than automatically true, and that important statistical/methodological claims receive a higher verification threshold.  NPTEL_Research_Methodology_Mast... +1

0. SOURCE / SCOPE

Week: 6
Module: 6
Lectures covered: 31–36
Lecture titles:

Lecture	Title
31	The Central Limit Theorem and its Applications — Part 01
32	The Central Limit Theorem and its Applications — Part 02
33	Error Bars and Confidence Interval — Part 01
34	Error Bars and Confidence Interval — Part 02
35	Measurement of a Proportion — Part 01
36	Measurement of a Proportion — Part 02

The transcript is reasonably understandable, although several mathematical expressions and symbols are corrupted by speech-to-text. The intended mathematics can generally be reconstructed from context.

External verification: Yes — especially for:

- the Central Limit Theorem,
- confidence intervals,
- Student's *t*-distribution,
- interpretation of confidence levels,
- binomial/proportion inference.

NIST confirms the standard confidence-interval formulation and, importantly, explicitly warns against interpreting a 95% confidence interval as a 95% probability that the particular interval contains the fixed population parameter. 

Important source limitation: The lecturer repeatedly uses “25 observations” as a **practical threshold for the CLT**. That is a teaching rule-of-thumb, not a universal theorem. The required sample size depends on the underlying distribution, dependence structure, tail behavior, desired approximation quality, and inferential goal.

1. LECTURE MAP

Lecture 31 — Central Limit Theorem and its Applications — Part 01

Main topic

The sampling distribution of the sample mean.

Central question

If I cannot measure an entire population, what can I infer about its mean from a finite sample?

Major concepts

- Population mean μ
- Population standard deviation σ
- Sample mean $\bar{x}$
- Sample standard deviation s
- Sampling
- Sampling distribution
- Central Limit Theorem
- Mean of sampling distribution
- Standard error
- Effect of sample size

Core result

For sufficiently large n , under the usual CLT conditions:

$$\bar{X} \approx N\left(\mu, \frac{\sigma^2}{n}\right)$$

or equivalently:

$$E[\bar{X}] = \mu$$

and

$$SD(\bar{X}) = \frac{\sigma}{\sqrt{n}}.$$

The transcript explicitly develops this idea: repeated samples of size n generate different means, and the distribution of those means approaches a normal distribution with mean μ and standard deviation $\sigma/\sqrt{n}$.  Week 6 research methodology tra...

Lecture 32 — Central Limit Theorem and its Applications — Part 02

Main topic

Applying the CLT to probability calculations.

Major concepts

- What "large n " means
- Sampling distribution
- z -score
- Standard error
- Probability from normal distribution
- Increasing n
- Estimating the population mean
- Confidence reasoning

Worked examples

The lecturer considers:

- population mean = 2,
- population SD = 0.7,
- sample sizes such as 20 and 50.

He calculates how the probability of obtaining a sample mean substantially above the population mean decreases as n increases.  Week 6 research methodology tra... +1

Lecture 33 — Error Bars and Confidence Interval — Part 01**Main topic**

How to communicate measurement uncertainty.

Major concepts

- Error bars
- Standard error
- Confidence level
- 68.3%
- 95%
- 99%
- z -values
- Reporting measurements
- Percentage error
- Significant digits / measurement precision
- 5σ in particle physics

The lecturer moves from:

sample $\rightarrow \bar{x}, s$

to:

estimated uncertainty around $\bar{x}$.

Lecture 34 — Error Bars and Confidence Interval — Part 02

Main topic

What happens when the sample is too small for the normal approximation?

Major concepts

- Error-bar width
- Sample-size effect
- $1/\sqrt{n}$
- Small samples
- Student's t -distribution
- Degrees of freedom
- Significance level
- Hypothesis testing
- One-sided evidence
- Example with $n = 9$

This lecture is particularly important because the course begins transitioning from descriptive measurement into **formal statistical inference**.

Lecture 35 — Measurement of a Proportion — Part 01

Main topic

Not every scientific measurement produces a continuous number.

Sometimes the outcome is:

0 or 1

and the quantity of interest is the **proportion of 1s**.

Examples

- tall vs short pea plants,
- brown vs white moths,
- two biological populations,
- evolutionary change.

Major concepts

- Binary observations
- Bernoulli variable
- Population proportion p
- Mean of Bernoulli variable
- Variance of Bernoulli variable
- Binomial distribution
- Counting successes
- Combinations

The lecturer develops the moth example with:

$$p = 0.6.$$

Lecture 36 — Measurement of a Proportion — Part 02

Main topic

Sampling distribution of an observed proportion.

Major concepts

- Number of successes X
- Measured proportion $\hat{p}$
- Binomial distribution
- Mean of $\hat{p}$
- Variance of $\hat{p}$
- Standard error of a proportion
- Approximate normality
- Rule of thumb $np > 10$ and $n(1 - p) > 10$
- Example with 33 successes out of 50

The central result is:

$$E[\hat{p}] = p$$

and approximately:

$$SD(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}.$$

The transcript derives this through the relationship between X , the number of 1s, and $\hat{p} = X/n$. [📄 Week 6 research methodology tra...](#)

2. EXECUTIVE UNDERSTANDING

What I Should Understand After This Week

1. We almost never observe an entire population.

We observe a **sample** and use it to infer properties of the underlying population.

2. A sample mean is itself a random quantity.

If you repeatedly draw samples of size n :

$$\bar{x}_1, \bar{x}_2, \bar{x}_3, \dots$$

they will differ.

Those means have their own distribution:

the sampling distribution of the mean.

3. The Central Limit Theorem explains why normal distributions appear so often in statistical inference.

For sufficiently large samples, the sampling distribution of the mean approaches a normal distribution under appropriate conditions.

Its spread is:

$$\frac{\sigma}{\sqrt{n}}.$$

4. Increasing sample size makes the estimate of the mean more precise.

Because:

$$SE(\bar{X}) = \frac{\sigma}{\sqrt{n}}.$$

Thus:

$$n \uparrow \Rightarrow SE \downarrow.$$

But the improvement follows a **square-root law**, not a linear law.

5. An error bar is meaningless unless you know what it represents.

An error bar might represent:

- SD,
- SE,
- confidence interval,
- measurement uncertainty,
- something else.

Therefore:

Never interpret an error bar without checking its definition.

6. Confidence intervals are inferential procedures, not probability statements about a fixed parameter.

This is one place where the lecturer's language needs refinement.

A 95% confidence procedure produces intervals that cover the fixed parameter in approximately 95% of repeated samples under the assumptions of the procedure. NIST explicitly makes this distinction.  NIST

7. Small samples require different methods.

When σ is unknown and the sample is small, the normal z -distribution is generally replaced by the **Student's t -distribution** under the standard normal-population

assumptions.

8. Continuous measurements and proportion measurements require different statistical models.

A measurement such as:

$$x = 17.3$$

is different from:

$$Y \in \{0, 1\}.$$

The latter leads naturally to:

Bernoulli $\rightarrow$ Binomial $\rightarrow$ Proportion inference.

9. The observed proportion $\hat{p}$ is itself random.

Different samples give different:

$$\hat{p}.$$

Its approximate standard error is:

$$SE(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}.$$

In practice, because p is usually unknown, it is commonly estimated using $\hat{p}$.

10. Statistical inference is essentially controlled reasoning under sampling uncertainty.

The whole week can be seen as:



3. THE ONE BIG IDEA

A finite sample does not tell us the population directly; it gives us a probability distribution over what sample-derived estimates could look like.

This is the conceptual leap of Week 6.

Before this week, you had:

sample $\rightarrow$ mean.

Now the course asks:

How much should I trust that mean?

The answer comes from studying the **distribution of means obtained from repeated samples**.

That leads naturally to:

Sampling distribution → standard error → confidence interval → hypothesis



And then the same logic is extended from numerical measurements to **proportions**.

4. DEEP CONCEPTUAL NOTES

4.1 Population vs Sample

[LECTURE]

The lecturer begins with an important distinction:

There is some population-level quantity "out there":

$$\mu, \sigma$$

but we only have:

$$\bar{x}, s.$$

For example:

Biology

Population:

all individuals of a species.

Target:

$$\mu = \text{true population average body mass.}$$

Sample:

25 individuals.

Observed:

$$\bar{x} = \text{sample average.}$$

Simple intuition

Imagine the true population average is hidden.

You cannot inspect the entire population.

You take:

10, 20, 50, 100

samples and calculate their means.

Each sample gives a slightly different answer.

The key question becomes:

How variable are those answers?

That variability is the **sampling distribution**.

4.2 Sampling Distribution

This is arguably the most important new concept.

Suppose the population is:

$$X \sim \text{some distribution.}$$

Take a sample of size n :

$$X_1, X_2, \dots, X_n.$$

Calculate:

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i.$$

Now repeat the entire sampling process many times.

You obtain:

$$\bar{X}_1, \bar{X}_2, \bar{X}_3, \dots$$

The distribution of these values is:

the sampling distribution of the sample mean.

Why it matters

This is what allows us to quantify:

"How unusual is the mean I observed?"

Without a sampling distribution, a number such as:

$$\bar{x} = 18.87$$

doesn't tell us much.

Once we understand its sampling distribution, we can ask:

How likely would a value this extreme be if the population mean were actually 21?

That is precisely the reasoning used in Lecture 34.

4.3 Central Limit Theorem

[LECTURE]

The lecturer states that when the sample size becomes sufficiently large:

the distribution of sample means approaches a normal distribution.

Its mean is:

$$\mu$$

and its standard deviation is:

$$\frac{\sigma}{\sqrt{n}}.$$

The transcript explicitly states this formulation.  Week 6 research methodology tra...

Precise formulation

Under standard independent-identically-distributed conditions, if:

$$E[X] = \mu$$

and

$$Var(X) = \sigma^2 < \infty,$$

then:

$$\boxed{\frac{\sqrt{n}(\bar{X} - \mu)}{\sigma} \xrightarrow{d} N(0, 1)}$$

as:

$$n \rightarrow \infty.$$

Equivalently, for sufficiently large n :

$$\boxed{\bar{X} \approx N\left(\mu, \frac{\sigma^2}{n}\right)}$$

This is the mathematical heart of the week's first two lectures.

4.4 Why Does the Standard Deviation Become $\sigma/\sqrt{n}$?

This is worth actually understanding.

Assume independent observations.

$$\bar{X} = \frac{X_1 + \cdots + X_n}{n}.$$

Then:

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{X_1 + \cdots + X_n}{n}\right).$$

For independent observations:

$$\text{Var}(X_1 + \cdots + X_n) = n\sigma^2.$$

Therefore:

$$\text{Var}(\bar{X}) = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n}.$$

Taking square root:

$$\boxed{SD(\bar{X}) = \frac{\sigma}{\sqrt{n}}}$$

This isn't merely a consequence of the CLT.

It follows from the variance properties of independent observations.

The CLT **provides the approximate normal shape.**

That distinction is important.

4.5 CLT Does Not Say "The Population Becomes Normal"

This is a very common misunderstanding.

Suppose:

$$X$$

is extremely skewed.

The CLT does **not** say:

the population distribution becomes normal.

It says:

the distribution of the sample mean becomes approximately normal as sample size increases.

This distinction is absolutely fundamental.

The transcript itself explicitly acknowledges that the population distribution can be non-normal.  Week 6 research methodology tra...

4.6 The "25 Samples" Claim

[LECTURE]

The lecturer repeatedly says approximately:

$n \geq 25$ is sufficient / "large."

He even calls 25 the answer to the question:

"How many readings do we need?"

The transcript states this explicitly.  Week 6 research methodology tra...

[VERIFIED / IMPORTANT QUALIFICATION]

This should **not** be memorised as:

$$n \geq 25 \Rightarrow \text{CLT always works}$$

That is not a universal theorem.

There is no single universal sample size at which the CLT suddenly becomes valid.

How quickly the sampling distribution approaches normality depends on the underlying distribution.

For example:

- nearly normal population $\rightarrow$ relatively small n may suffice;
- moderately skewed population $\rightarrow$ larger n ;
- extremely skewed/heavy-tailed population $\rightarrow$ potentially much larger n ;
- infinite-variance distributions can require special treatment because the classical CLT assumptions fail.

Correct mental model

25 is a course-level rule of thumb, not a mathematical boundary.

This is one of the most important corrections in Week 6.

4.7 Standard Error

The lecturer introduces:

$$SE(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

and in practice estimates it using:

$$\widehat{SE}(\bar{X}) = \frac{s}{\sqrt{n}}.$$

The transcript explicitly identifies the standard deviation of the sampling distribution of the mean as the "standard error of the mean."  Week 6 research methodology tra...

Intuition

Standard deviation asks:

How much do individual observations vary?

Standard error asks:

How much would the **sample mean** vary if I repeated the sampling process?

That distinction is crucial.

Example

Suppose:

$$s = 10$$

and:

$$n = 100.$$

Then:

$$SE = \frac{10}{\sqrt{100}} = 1.$$

If:

$$n = 400,$$

then:

$$SE = \frac{10}{20} = 0.5.$$

So quadrupling n halves the standard error.

4.8 The Square-Root Law

This becomes one of the most useful experimental-design rules.

$$SE \propto \frac{1}{\sqrt{n}}.$$

Therefore:

Double sample size

$$n \rightarrow 2n$$

gives:

$$SE \rightarrow \frac{SE}{\sqrt{2}}.$$

Four times sample size

$$n \rightarrow 4n$$

gives:

$$SE \rightarrow \frac{SE}{2}.$$

The lecturer explicitly explains that halving an error bar requires approximately four times as many observations.  Week 6 research methodology tra...

This is a very important experimental-design insight.

4.9 Probability Example from Lecture 32

The lecturer uses:

$$\mu = 2$$

$$\sigma = 0.7$$

and:

$$n = 20.$$

Then:

$$SE = \frac{0.7}{\sqrt{20}} \approx 0.1565.$$

Suppose we ask:

$$P(\bar{X} > 2.2).$$

The corresponding standardized value is:

$$z = \frac{2.2 - 2}{0.1565} \approx 1.28.$$

Using the standard normal distribution:

$$P(Z < 1.28) \approx 0.8997.$$

Therefore:

$$P(Z > 1.28) = 1 - 0.8997 = 0.1003.$$

So approximately:

$$P(\bar{X} > 2.2) \approx 10\%.$$

The transcript works through this calculation.  Week 6 research methodology tra...

What changes if $n = 50$?

$$SE = \frac{0.7}{\sqrt{50}} \approx 0.099.$$

Then:

$$z = \frac{2.2 - 2}{0.099} \approx 2.02.$$

Thus:

$$P(\bar{X} > 2.2) \approx 2.2\%.$$

So larger n makes large deviations of the sample mean less probable.

This is one of the practical powers of the CLT.

4.10 Confidence Intervals

The course now moves from:

"What is the probability of observing this sample mean?"

to:

"What range of population means is reasonably compatible with our measurement?"

Conceptually:

$$\text{estimate} \pm \text{uncertainty}$$

For known σ under normality:

$$\bar{x} \pm z_{\alpha/2} \frac{\sigma}{\sqrt{n}}$$

If σ is unknown and the normal-population assumptions are appropriate:

$$\bar{x} \pm t_{\alpha/2, n-1} \frac{s}{\sqrt{n}}$$

NIST gives this standard t -based confidence interval formulation. NIST

4.11 68.3%, 95%, 99%

For a standard normal distribution:

Approximately 68.3%

$$\mu \pm 1\sigma.$$

Approximately 95%

$$\mu \pm 1.96\sigma.$$

Approximately 99%

$$\mu \pm 2.576\sigma.$$

The lecturer uses:

$$2.58$$

for 99%, which is appropriate as a rounded value. The transcript applies these factors to the **standard error of the mean** when discussing uncertainty in the sample mean.

 Week 6 research methodology tra...

4.12 A Critical Correction: Confidence Is Not Probability of the Parameter

This is perhaps the **single most important statistical correction** in Week 6.

The lecturer repeatedly phrases statements approximately as:

"There is a 95% probability that the actual mean lies within this interval."

That wording is common in informal teaching, but it is **not the standard frequentist interpretation**.

Correct frequentist interpretation

Suppose a procedure generates:

$$CI(X)$$

with 95% coverage.

Then:

If we repeatedly drew samples and constructed intervals using this same procedure, approximately 95% of those intervals would contain the fixed population parameter.

NIST explicitly states this distinction. NIST NIST

After you calculate one particular interval:

$$[10.2, 11.7],$$

the fixed parameter either lies inside it or does not.

The confidence level describes the **long-run performance of the procedure**, not a posterior probability assigned to that particular interval.

4.13 Error Bars

The lecturer explains error bars as graphical representations of uncertainty.

For example:

$$\bar{x} \pm SE.$$

But this raises a major research-practice point:

An error bar is not self-interpreting.

Different papers use error bars for:

- SD,
- SE,
- 95% CI,
- measurement uncertainty,
- prediction intervals,
- other quantities.

Therefore a rigorous paper must specify what the error bars represent.

4.14 Why Error Bars Matter

Suppose two experiments report:

Experiment A

$$10.0 \pm 0.1$$

Experiment B

$$10.2 \pm 1.5$$

The point estimates alone suggest:

$$10.0 \quad \text{vs} \quad 10.2.$$

But the uncertainty tells us how precise those estimates are.

A researcher therefore shouldn't merely ask:

"Which mean is larger?"

but:

"How uncertain are these estimates, and what inferential claim does that uncertainty support?"

4.15 Reporting Significant Digits

The lecturer makes a very practical point:

Do not report more precision than your measurement apparatus justifies.

For example, if an instrument only meaningfully resolves:

$$0.001$$

units, reporting:

3.14159265358979

is misleading.

The transcript explicitly discusses the instrument's least count and the need to avoid meaningless extra decimal places.  Week 6 research methodology tra...

Important nuance

The exact number of reported digits depends on:

- instrument resolution,
- calibration,
- uncertainty,
- measurement model,
- significant-figure conventions.

It isn't determined solely by the displayed digits of the calculator.

4.16 5σ in Particle Physics

The lecturer explains the famous:

$$5\sigma$$

threshold.

For a standard normal distribution, the **one-sided** tail beyond 5σ is approximately:

$$2.87 \times 10^{-7},$$

or about:

$$1/3.5 \text{ million.}$$

For a **two-sided** tail, the probability is approximately:

$$5.73 \times 10^{-7}.$$

This distinction matters because the transcript's "one in 3.5 million" corresponds approximately to the **one-sided** tail.

Important research qualification

The lecturer's wording makes 5σ sound like:

"probability the discovery is false = 1 in 3.5 million."

That is **not correct**.

A 5σ result corresponds to a very small **p-value under a specified null model**, not the probability that the null hypothesis is true or that the discovery is false.

This is an important distinction between:

$$P(\text{data} \mid H_0)$$

and:

$$P(H_0 \mid \text{data}).$$

They are not the same.

4.17 Small Samples and Student's t

Lecture 34 asks:

What happens when n is small?

If the population standard deviation σ is unknown, we estimate it with:

$$s.$$

That introduces additional uncertainty.

For normal-population data, the appropriate standardized statistic becomes:

$$t = \frac{\bar{X} - \mu_0}{s/\sqrt{n}}$$

and follows a Student's t -distribution with:

$$df = n - 1$$

under the null hypothesis and the usual assumptions.

NIST confirms the t -based formulation and $n - 1$ degrees of freedom. NIST NIST

4.18 Why Does the t -Distribution Have Heavier Tails?

Because:

$$\sigma$$

is unknown.

We replace it with:

$$s.$$

That introduces additional uncertainty.

Therefore the t -distribution has heavier tails than the standard normal distribution, especially for small n .

As:

$$n \rightarrow \infty,$$

the t -distribution approaches:

$$N(0, 1).$$

Thus:

t is especially important for small samples and becomes increasingly similar to z as sample size grows.

4.19 Lecture 34 Example

The lecturer has:

$$n = 9$$

measurements of particle mass:

$$16.2, 19.7, 21.8, 15.6, 19.0, 18.7, 16.9, 21.7, 20.2.$$

He obtains approximately:

$$\bar{x} = 18.87$$

and:

$$s \approx 2.2583.$$

The question:

Is there sufficient evidence that the true mean mass is below 21?

The lecturer defines "sufficient evidence" as requiring:

$$P(\text{wrong conclusion}) < 1\%.$$

Then he considers:

$$H_0 : \mu = 21$$

and calculates:

$$t = \frac{18.87 - 21}{2.2583/\sqrt{9}} \approx -2.83.$$

Degrees of freedom:

$$df = 9 - 1 = 8.$$

At a one-sided 1% significance level, the critical magnitude is approximately:

$$3.355.$$

Since:

$$|-2.83| < 3.355,$$

the evidence is insufficient to reject the null at the 1% level in favor of:

$$\mu < 21.$$

This matches the lecturer's conclusion.  Week 6 research methodology tra...

4.20 "Cannot Reject" Is Not "Proved Equal"

This is a crucial statistical habit.

If the result isn't sufficiently extreme:

Fail to reject H_0 .

Do not say:

"We proved H_0 is true."

In the particle example, the correct conclusion is:

There is insufficient evidence at the specified significance level to conclude that $\mu < 21$.

Not:

"The particle definitely has mean mass 21."

4.21 Measurement of a Proportion

Now the course changes the type of variable.

Instead of:

$$X = 17.82$$

we have:

$$Y \in \{0, 1\}.$$

Examples:

- tall/short,
- brown/white,
- diseased/non-diseased,
- success/failure.

The scientific quantity of interest is often:

$$p = P(Y = 1).$$

4.22 Bernoulli Variable

Define:

$$Y = \begin{cases} 1 & \text{success/event of interest} \\ 0 & \text{otherwise} \end{cases}$$

with:

$$P(Y = 1) = p$$

and:

$$P(Y = 0) = 1 - p.$$

Mean

$$E[Y] = 1(p) + 0(1 - p)$$

so:

$$E[Y] = p.$$

Variance

$$Var(Y) = p(1 - p).$$

Therefore:

$$SD(Y) = \sqrt{p(1 - p)}.$$

The lecturer derives exactly this using the brown/white moth example.

 Week 6 research methodology tra...

4.23 Why Is the Mean of a 0/1 Variable a Proportion?

This is a beautiful and useful connection.

Suppose:

$$Y_1, Y_2, \dots, Y_n$$

are binary observations.

Then:

$$Y_1 + \dots + Y_n = X$$

where X is the number of successes.

The sample mean is:

$$\bar{Y} = \frac{Y_1 + \dots + Y_n}{n} = \frac{X}{n}.$$

But:

$$\frac{X}{n}$$

is exactly the observed proportion.

Therefore:

$$\hat{p} = \bar{Y}.$$

This means:

A proportion is simply the sample mean of a 0/1 variable.

That is a very powerful statistical connection.

4.24 Binomial Distribution

If we have:

$$n$$

independent Bernoulli trials with success probability p , then the number of successes:

$$X$$

follows:

$$X \sim \text{Binomial}(n, p).$$

The probability of exactly k successes is:

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}.$$

The lecturer derives this through the moth example.  Week 6 research methodology tra...

4.25 Why Does the Combination Factor Appear?

Suppose:

$$n = 10, \quad k = 6.$$

One particular sequence might be:

$$0011111100.$$

Its probability is:

$$(0.4)^4 (0.6)^6.$$

But there are many possible arrangements of six 1s among ten positions.

The number is:

$$\binom{10}{6}.$$

Therefore:

$$P(X = 6) = \binom{10}{6} (0.6)^6 (0.4)^4.$$

The lecturer calculates approximately:

$$0.2508.$$

That is:

even though the population proportion is 60%, a particular sample does not have to contain exactly 60% successes.

This is a very important sampling concept.

4.26 Sampling Distribution of the Proportion

Define:

$$\hat{p} = \frac{X}{n}.$$

Since:

$$E[X] = np$$

we get:

$$E[\hat{p}] = E\left[\frac{X}{n}\right] = \frac{np}{n} = p.$$

Therefore:

$$\boxed{E[\hat{p}] = p}.$$

The observed proportion is an **unbiased estimator** of p under the standard binomial model.

4.27 Standard Deviation of the Proportion

Since:

$$\text{Var}(X) = np(1 - p),$$

and:

$$\hat{p} = \frac{X}{n},$$

then:

$$\begin{aligned}
 \text{Var}(\hat{p}) &= \frac{1}{n^2} \text{Var}(X) \\
 &= \frac{np(1-p)}{n^2} \\
 &= \frac{p(1-p)}{n}.
 \end{aligned}$$

Therefore:

$$SD(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}$$

This is the standard error of the observed proportion under the binomial model.

The transcript derives this explicitly. [📄 Week 6 research methodology tra...](#)

4.28 Normal Approximation to the Binomial

The binomial distribution is discrete:

$$X = 0, 1, 2, \dots, n.$$

For sufficiently large expected counts, its shape can be approximated by a normal distribution.

The lecturer gives the rule of thumb:

$$np > 10$$

and:

$$n(1-p) > 10.$$

He explicitly says this is a **rule of thumb, not a theorem**. [📄 Week 6 research methodology tra...](#)

That qualification is correct and important.

4.29 Practical Problem: p Is Unknown

There is a circular-looking problem.

We want:

$$SE(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}$$

but we don't know:

$$p.$$

So we commonly substitute:

$$\hat{p}$$

for:

$$p.$$

giving the estimated standard error:

$$\widehat{SE}(\hat{p}) = \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}.$$

The lecturer uses precisely this substitution in Lecture 36.  Week 6 research methodology tra...

4.30 Lecture 36 Example

Suppose:

$$n = 50$$

and:

$$X = 33.$$

Then:

$$\hat{p} = \frac{33}{50} = 0.66.$$

Estimated standard error:

$$\begin{aligned} & \sqrt{\frac{0.66(1 - 0.66)}{50}} \\ &= \sqrt{\frac{0.2244}{50}} \\ &\approx 0.067. \end{aligned}$$

The lecturer then asks whether the population proportion could plausibly lie within:

$$0.64 \leq p \leq 0.68.$$

The interval has half-width:

$$0.02.$$

Standardised:

$$z \approx \frac{0.02}{0.067} \approx 0.30.$$

The corresponding central normal probability is only about:

$$23.6\%.$$

The lecturer concludes that this range has low confidence.  Week 6 research methodology tra...

5.31 A Subtle but Important Issue With the Proportion Example

The lecturer's arithmetic is basically consistent, but his interpretation again uses probability language in the frequentist way that should be refined.

It is better to say:

Under the assumed approximate sampling distribution, an interval constructed as 0.66 ± 0.02 has a coverage probability corresponding to roughly 23.6% **if the procedure and normal approximation are appropriate.**

It is not best phrased as:

"There is a 23.6% probability that the true p lies in this particular interval."

That is the same confidence-interval interpretation issue encountered earlier.

5.32 Why the Course Moves From Means to Proportions

This is an important structural point.

The lecturer is showing that scientific measurement isn't limited to:

continuous physical quantities.

We may measure:

- mass,
- length,
- temperature,

but also:

- proportion infected,
- proportion surviving,
- proportion exhibiting a trait,
- proportion choosing one option,
- proportion of defective products.

The underlying methodology remains:

Population → Sample → Sampling distribution → Uncertainty → Inference.

Only the statistical model changes.

5.33 Measurement Theory → Statistical Inference

The course's progression is now very clear:

Weeks 5–6

Measurement

↓

Repeated observations

↓

Sample statistics

↓

Sampling distributions

↓

Uncertainty

↓

Confidence intervals

↓

Hypothesis tests

This is the statistical machinery needed to turn raw measurements into defensible scientific claims.

6. LECTURER CLAIM AUDIT

Lecturer's claim	Verdict	Assessment	Qualification / correction	Evidence
Sample means approach a normal distribution for sufficiently large n	✓ Correct	Core CLT result under suitable assumptions	"Sufficiently large" depends on the underlying distribution and conditions	NIST/statistical t 📄 Week 6 research
SD of sample mean is $\sigma/\sqrt{n}$	✓ Correct	Correct for independent observations with finite variance	Dependence changes the variance formula	Lecture 📄 Week 6 research
$n = 25$ is the required sample size	● Potentially misleading	Useful course-level rule of thumb	No universal CLT cutoff exists	Lecture explicitly practical threshc 📄 Week 6 research
Increasing n makes the sampling distribution narrower	✓ Correct	Directly follows from $\sigma/\sqrt{n}$	Assuming comparable sampling conditions	Lecture 📄 Week 6 research
Standard error is SD divided by $\sqrt{n}$	● Correct but simplified	Correct for the mean under the standard independent-sampling model	More generally, SE depends on estimator and sampling design	Lecture 📄 Week 6 research

Lecturer's claim	Verdict	Assessment	Qualification / correction	Evidence
An error bar represents confidence that the actual value lies in the bar	 Potentially misleading	Depends on what the error bar represents	Error bars may represent SD, SE, CI, etc.; confidence interpretation must be specified	Lecture  Week 6 research
A 95% interval means 95% probability the fixed mean is inside it	 Potentially misleading	Informal wording, but not standard frequentist interpretation	95% refers to long-run coverage of the interval-producing procedure	NIST  NIST
5σ means about 1 in 3.5 million random outcomes	 Correct but simplified	Approximately correct for a one-sided standard-normal tail	Does not equal probability that the null is true or discovery is false	Statistical interpi
5σ means only a 1-in-3.5-million probability the discovery is false	 Incorrect	Confuses $P(\text{data} H_0)$ with $P(H_0 \text{data})$	A p-value is not posterior probability of the null	Statistical inferen
If $n < 25$, normal methods cannot be used	 Potentially misleading	Too absolute	Depends on assumptions and distribution; t -methods can be valid under normal-population assumptions for small n	NIST  NIST
Small samples use t -distribution	 Correct but simplified	Appropriate for inference about a normal population mean when σ is unknown	Not simply "small n = automatically use t " in every setting	NIST  NIST
Degrees of freedom for one-sample t is $n - 1$	 Correct	Standard one-sample case	Depends on test/model	NIST  NIST
Binary measurement can be represented as 0/1	 Correct	Standard Bernoulli representation	Coding convention is arbitrary; 1/0 assignment can be reversed	Lecture  Week 6 research
Mean of Bernoulli variable is p	 Correct	Fundamental result	Assumes $Y \in \{0, 1\}$ with $P(Y = 1) = p$	Lecture  Week 6 research
Variance of Bernoulli is $p(1 - p)$	 Correct	Fundamental	—	Lecture  Week 6 research
$np > 10$ and $n(1 - p) > 10$ guarantee normality	 Context-dependent	Useful rule of thumb	Approximation quality depends on purpose; not a theorem or guarantee	Lecturer explicitl of thumb  Week 6 research

Lecturer's claim	Verdict	Assessment	Qualification / correction	Evidence
Measured proportion has mean p	✓ Correct	$\hat{p}$ is unbiased under the binomial model	Requires the assumed sampling model	Lecture 📄 Week 6 research
Standard deviation of $\hat{p}$ is $\sqrt{p(1-p)/n}$	✓ Correct	Exact under independent binomial sampling	Unknown p must be estimated in practice	Lecture 📄 Week 6 research

7. BEYOND THE LECTURE

7.1 The CLT Requires Conditions

The course presents the CLT intuitively.

The more precise researcher-level picture is that the classical CLT requires assumptions such as:

- independent or suitably weakly dependent observations,
- finite variance,
- appropriate identically distributed or suitably controlled distributions, depending on the version of the theorem.

This matters because real scientific data may contain:

- time dependence,
- spatial correlation,
- repeated measurements,
- clustered observations.

For example, 100 measurements from the same sensor taken one millisecond apart are **not automatically equivalent to 100 independent measurements**.

7.2 Effective Sample Size

Suppose:

$$X_1, X_2, \dots, X_n$$

are strongly correlated.

Then:

$$\text{Var}(\bar{X})$$

can be substantially larger than:

$$\frac{\sigma^2}{n}.$$

Therefore the naïve formula:

$$SE = \frac{s}{\sqrt{n}}$$

can underestimate uncertainty.

This is highly relevant to:

- time series,
- longitudinal studies,
- spatial measurements,
- MCMC,
- reinforcement learning,
- repeated evaluation of AI systems.

7.3 Confidence Intervals Are Procedures

This deserves permanent memory.

A confidence interval isn't:

"the probability range where the parameter lives."

It is:

an output of a procedure whose repeated-sampling coverage is designed to be approximately $1 - \alpha$.

NIST explicitly explains this point. NIST NIST

7.4 There Is More Than One Way to Build a Confidence Interval

For a population mean:

- z -interval,
- t -interval,
- bootstrap interval,
- Bayesian credible interval — conceptually different framework.

For a proportion:

- Wald interval,
- Wilson interval,
- exact/binomial intervals,
- score-based intervals.

The course's:

$$\hat{p} \pm z \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

is an introductory approximation and can behave poorly, especially with small samples or proportions near 0 or 1.

NIST provides exact binomial confidence-interval methods that are not simply symmetric around $\hat{p}$. NIST NIST

7.5 Why the Simple Proportion Interval Can Fail

Suppose:

$$n = 10$$

and:

$$X = 0.$$

Then:

$$\hat{p} = 0.$$

The naïve standard error is:

$$\sqrt{\frac{0(1-0)}{10}} = 0.$$

That would imply:

zero uncertainty.

Obviously that is not a sensible inferential conclusion.

This is one reason more robust proportion intervals such as **Wilson** or exact binomial methods are often preferred over the simple Wald interval.

7.6 Statistical Significance Is Not Scientific Importance

Suppose an AI system improves:

$$90.00\% \rightarrow 90.05\%.$$

With an enormous test set, the difference could be statistically detectable.

But:

Is a 0.05 percentage-point improvement scientifically or practically important?

That is a different question.

Conversely, a genuinely important effect may fail to reach conventional statistical significance in a small study.

Thus:

statistical significance $\neq$ scientific importance

8. RESEARCHER'S PERSPECTIVE

What a Good Researcher Should Notice

1. "How many samples?" is not a universal number.

The answer depends on:

- variance,
 - desired precision,
 - confidence level,
 - effect size,
 - statistical power,
 - sampling design,
 - cost.
-

2. More data are not automatically independent information.

Ten thousand correlated measurements may contain less information than one thousand approximately independent observations.

3. The error bar must have a definition.

When reading a paper:

First question: what exactly does the error bar represent?

4. Uncertainty should be planned before the experiment.

Don't collect arbitrary data and only afterward ask:

"Was this enough?"

Ideally determine:

- target precision,
- acceptable error,
- confidence level,
- statistical power,
- sample size.

before data collection.

5. "No significant difference" is not the same as "identical."

It may simply mean:

evidence was insufficient to detect the proposed difference.

6. A p-value is not the probability that your hypothesis is true.

This distinction will become increasingly important as the course develops.

7. Statistical assumptions are part of the scientific argument.

When you use:

t -test

or:

normal approximation,

you are making assumptions.

Those assumptions need to be justified.

9. CONNECTION TO AI/ML RESEARCH

This week's material is **extremely relevant to AI research**, especially experimental evaluation.

9.1 Model Performance Is a Sample Statistic

Suppose you evaluate an LLM on:

$$n = 1,000$$

test questions.

You observe:

$$\hat{p} = 0.84.$$

That 84% is not necessarily:

"the model's true universal accuracy."

It is an estimate based on a sample of questions.

Conceptually:

Benchmark accuracy = sample statistic

and your actual research question may concern some broader target population of tasks.

9.2 Benchmark = Sampling Process

Suppose your intended claim is:

"Model A performs better on scientific question answering."

You need to ask:

What is the population?

All possible scientific questions?

Questions from a particular benchmark?

Questions in a particular domain?

English questions only?

Graduate-level questions?

What is the sample?

Your 500 evaluation questions.

Then:

$$\hat{p}_A$$

and:

$$\hat{p}_B$$

are sample statistics.

The difference:

$$\hat{p}_A - \hat{p}_B$$

also has uncertainty.

9.3 Example: Comparing Two AI Systems

Suppose:

Model A

$$840/1000 = 84\%.$$

Model B

$$860/1000 = 86\%.$$

A naïve conclusion:

B is better.

A researcher asks:

Is the 2-point difference larger than the uncertainty associated with the evaluation?

Then consider:

- paired vs independent questions,

- sampling variability,
- bootstrap confidence intervals,
- appropriate statistical tests,
- multiple comparisons,
- benchmark contamination,
- question dependence.

This is exactly the methodological mindset Week 6 is developing.

9.4 Why Test Examples Are Not Always Independent

This is especially important in AI.

Suppose a benchmark contains:

- 100 questions generated from one source,
- 50 paraphrases of the same question,
- multiple questions testing the same underlying fact.

Then:

$$n = 1000$$

does not necessarily mean:

1000 independent pieces of evidence.

If observations are correlated, naïvely applying:

$$SE = \frac{s}{\sqrt{n}}$$

can produce overconfident conclusions.

9.5 Repeated Random Seeds

Suppose you train the same model five times:

91.2, 90.7, 91.6, 91.0, 91.4.

These give information about **run-to-run variability**.

But they do not necessarily tell you how the model will perform on a new population of tasks.

You must distinguish:

Variability across seeds

from:

Variability across evaluation examples

from:

Variability across datasets/distributions.

These are different sources of uncertainty.

9.6 AI Research Example — Your RAG System

Suppose you want to test:

H: Hybrid retrieval improves RAG answer quality compared with dense retrieval alone.

You evaluate:

$$n = 500$$

queries.

You obtain:

Dense

82.0%

Hybrid

84.4%.

Your first instinct might be:

"Hybrid improves accuracy by 2.4 points."

A research-grade analysis asks:

1. What population do the 500 queries represent?
2. Are queries independent?
3. Was the test set fixed before experimentation?
4. Did you tune the system on this test set?
5. How many configurations were tried?
6. Are the two systems evaluated on exactly the same questions?
7. Is the 2.4-point difference larger than sampling uncertainty?
8. Does the improvement survive different random seeds?
9. Does it generalize to another dataset?
10. What mechanism explains the improvement?

That is **Week 6 methodology applied directly to your AI research.**

9.7 AI Error Bars

Instead of reporting:

Accuracy = 84.4%

you might report:

Accuracy = 84.4% with an appropriate confidence interval.

But the interval should match the data-generating situation.

For binary correctness on independent evaluation examples, a binomial-based method may be appropriate.

For correlated examples:

a simple binomial interval may be inappropriate.

For repeated seeds:

the relevant unit of variation may be the training run.

For multiple datasets:

dataset-level variability may be the more meaningful source of uncertainty.

So:

There is no universal "AI error bar."

9.8 AI Connection — Sample Size and Benchmark Design

Suppose you want to detect a 1 percentage-point difference.

You cannot simply say:

"I'll use 25 test questions because the CLT says 25."

You need to consider:

- baseline accuracy,
- effect size,
- significance level,
- desired power,
- paired design,
- variance,
- dependence,
- multiple comparisons.

This is a major upgrade over blindly applying the lecturer's 25-observation rule.

9.9 AI Connection — Proportions

Many AI metrics are proportions.

Examples:

$$\text{Accuracy} = \frac{\text{correct predictions}}{\text{total predictions}}$$

$$\text{Success rate} = \frac{\text{successful runs}}{\text{total runs}}$$

$$\text{Tool-call success} = \frac{\text{successful tool calls}}{\text{total calls}}.$$

These naturally connect to the binomial framework.

10. RESEARCH EXAMPLES

Example 1 — Hypothetical AI Benchmark

Research problem

Does Model A outperform Model B?

Research question

Is the accuracy difference between A and B greater than expected from sampling variability?

Hypothesis

$$H_0 : p_A = p_B$$

versus:

$$H_1 : p_A \neq p_B.$$

Variables

- independent variable: model
- dependent variable: correctness
- sampling unit: evaluation question

Method

Evaluate both models on the same pre-specified test set.

Evidence

Accuracy difference + appropriate uncertainty/statistical test.

Possible confounders

- different prompts,
- different temperature,
- different context lengths,
- leakage,
- benchmark contamination.

Rigorous researcher

Would pre-specify:

- evaluation set,
- metric,
- comparison,
- analysis method.

Example 2 — Hypothetical RAG Ablation

Research question

Does RRF contribute to hybrid retrieval performance?

Hypothesis

H : RRF improves retrieval effectiveness.

Compare

- Dense
- BM25
- Dense + BM25 without RRF
- Dense + BM25 + RRF

Measurement

Recall@k, MRR, nDCG.

Statistical question

Are observed differences larger than expected from query-level sampling variation?

Alternative explanation

Perhaps the improvement is actually due to:

- increased retrieval depth,
- duplicate removal,
- changed ranking,
- more documents entering the context.

This is where the statistical methodology connects to experimental control.

Example 3 — Hypothetical Model Seed Experiment

Question

How stable is model performance across training runs?

Run:

$$n = 10$$

independent seeds.

Record:

$$x_1, \dots, x_{10}.$$

Compute:

$$\bar{x}, s.$$

Then report an uncertainty measure appropriate to the intended claim.

But:

This does **not** automatically justify treating the 10 seeds as the entire source of uncertainty in model generalisation.

Seed variability and dataset variability are distinct.

11. COMMON RESEARCHER MISTAKES

Mistake 1 — "The CLT says 25 is enough."

Why it happens

Memorising the lecture's rule.

Why problematic

25 is not universal.

Fix

Ask:

What is the underlying distribution and what precision do I need?

Mistake 2 — Confusing SD and SE

Error

$$SD = s$$

is not the same as:

$$SE = \frac{s}{\sqrt{n}}.$$

Consequence

You may dramatically misrepresent uncertainty.

Mistake 3 — Treating an error bar as self-explanatory

Fix

Always determine:

SD? SE? CI? Something else?

Mistake 4 — Saying "95% probability the true value is inside"

Fix

Use the correct confidence-interval interpretation.

Mistake 5 — Thinking more observations always solve everything

They reduce random uncertainty under appropriate assumptions.

They do not automatically eliminate:

- bias,
 - confounding,
 - dependence,
 - bad sampling.
-

Mistake 6 — Treating a non-significant result as proof of no effect**Fix**

Look at:

- effect estimate,
 - confidence interval,
 - power,
 - practical significance.
-

Mistake 7 — Applying a normal approximation blindly to proportions**Fix**

Check whether the approximation is appropriate and consider more robust intervals for small/extreme counts.

Mistake 8 — Treating correlated observations as independent

This is particularly dangerous in:

- time series,
 - biological repeated measures,
 - spatial data,
 - benchmark clusters,
 - multiple questions generated from the same source.
-

Mistake 9 — Reporting excessive decimal places**Fix**

Match reported precision to meaningful measurement uncertainty.

12. TERMINOLOGY / GLOSSARY

Population — the target set or underlying distribution about which inference is desired.

Sample — observed subset of the population.

Population parameter — fixed but usually unknown quantity such as μ , σ , or p .

Sample statistic — quantity calculated from observed data, such as $\bar{x}$, s , or $\hat{p}$.

Sampling distribution — distribution of a statistic across repeated samples.

Central Limit Theorem — result describing convergence of suitably standardized sums/means toward a normal distribution under appropriate conditions.

Standard error — standard deviation of a sampling distribution or estimated standard deviation of an estimator.

Confidence interval — interval produced by a procedure with specified long-run coverage.

Confidence level — nominal long-run coverage of the confidence-interval procedure.

Significance level α — threshold for rejecting a null hypothesis in a specified hypothesis-testing framework.

z -score — standardized value relative to a normal distribution.

t -statistic — standardized statistic using an estimated standard deviation, with a Student's t -distribution under the relevant assumptions.

Degrees of freedom — number of independent pieces of information available after accounting for estimated parameters/constraints.

Bernoulli variable — binary random variable taking 0 or 1.

Binomial variable — number of successes in n independent Bernoulli trials with common success probability p .

Population proportion p — probability/proportion of population members satisfying a specified binary condition.

Sample proportion $\hat{p}$ — observed fraction of successes.

Statistical significance — evidence against a specified null hypothesis at a chosen threshold; not synonymous with practical importance.

13. MATHEMATICS / TECHNICAL DETAILS

Sample mean

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

Sampling distribution of mean

$$E[\bar{X}] = \mu$$

and, for independent observations:

$$\text{Var}(\bar{X}) = \frac{\sigma^2}{n}.$$

Thus:

$$SE(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

or estimated:

$$\widehat{SE}(\bar{X}) = \frac{s}{\sqrt{n}}.$$

CLT standardisation

$$Z = \frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$$

approximately follows:

$$N(0, 1)$$

for sufficiently large n under appropriate assumptions.

Normal confidence interval

When σ is known:

$$\bar{x} \pm z_{\alpha/2} \frac{\sigma}{\sqrt{n}}$$

One-sample t -interval

When σ is unknown and the normal-population assumptions are appropriate:

$$\bar{x} \pm t_{\alpha/2, n-1} \frac{s}{\sqrt{n}}$$

NIST gives this standard formulation. NIST NIST

One-sample t -test

$$t = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$$

with:

$$df = n - 1.$$

Bernoulli mean

$$E[Y] = p$$

Bernoulli variance

$$\text{Var}(Y) = p(1 - p)$$

Binomial distribution

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

Binomial mean

$$E[X] = np$$

Binomial variance

$$\text{Var}(X) = np(1 - p)$$

Sample proportion

$$\hat{p} = \frac{X}{n}$$

Expected sample proportion

$$E[\hat{p}] = p$$

Variance of sample proportion

$$Var(\hat{p}) = \frac{p(1-p)}{n}$$

Standard error of proportion

$$SE(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}$$

Estimated using:

$$\widehat{SE}(\hat{p}) = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

14. CRITICAL THINKING QUESTIONS

Q1. Why does increasing n reduce the standard error only as $1/\sqrt{n}$?

Because:

$$Var(\bar{X}) = \frac{\sigma^2}{n},$$

so taking the square root gives:

$$SE = \frac{\sigma}{\sqrt{n}}.$$

Q2. Does the CLT say the population becomes normal?

No.

It concerns the **sampling distribution of an estimator**, especially the sample mean.

Q3. Is $n = 25$ a mathematical threshold for the CLT?

No.

It is a course-level rule of thumb.

Q4. Why is the sample mean more stable than individual observations?

Averaging independent observations causes independent fluctuations to partially cancel, reducing variance by a factor of n .

Q5. What is the difference between SD and SE?

SD:

spread of individual observations.

SE:

spread of the estimator across repeated samples.

Q6. Why does the t -distribution appear?

Because σ is unknown and replaced by s , introducing additional uncertainty.

Q7. Why does t have heavier tails than z ?

It accounts for uncertainty in estimating the population standard deviation.

Q8. What does "fail to reject H_0 " mean?

The observed evidence is not sufficiently inconsistent with H_0 under the chosen statistical criterion.

It does **not** prove H_0 .

Q9. Why can a proportion be treated as a mean?

Because a binary observation is 0 or 1:

$$\bar{Y} = \frac{\sum Y_i}{n} = \frac{\text{number of 1s}}{n} = \hat{p}.$$

Q10. Why does the binomial distribution appear?

Because X counts successes across n Bernoulli trials.

Q11. Why can't we blindly use the normal approximation when p is close to zero?

Because the distribution becomes strongly skewed/discrete and expected success counts may be too small.

Q12. Why isn't a 95% confidence interval a 95% probability statement about the parameter?

Because, in the classical frequentist framework, the parameter is treated as fixed; the random object is the interval-producing procedure. NIST NIST

Q13. Why doesn't four times the sample size always make your scientific conclusion four times better?

Because precision improves as:

$$1/\sqrt{n},$$

and systematic errors, bias and model misspecification are not automatically fixed.

Q14. Why can 10,000 observations still be statistically misleading?

If they are:

- strongly correlated,
- biased,
- nonrepresentative,
- contaminated,
- measured with systematic error.

15. ANSWERS / REASONING

The unifying answer is:

Statistical inference is only as good as the relationship between the observed sample and the population/process you are trying to infer about.

The formulas quantify sampling uncertainty **under assumptions**.

They don't magically validate those assumptions.

16. WHAT A GOOD RESEARCHER SHOULD NOTICE

1. The estimator is not the parameter.

Always distinguish:

$$\bar{x}$$

from:

$$\mu.$$

And:

$$\hat{p}$$

from:

$$p.$$

2. The sampling distribution is the hidden engine of inference.

Confidence intervals and hypothesis tests make sense because we know something about:

$$P(\text{statistic} \mid \text{population/model}).$$

3. Every statistical formula contains assumptions.

Before applying:

$$SE = \frac{s}{\sqrt{n}},$$

ask:

Are the observations sufficiently independent?

4. Sample size affects precision, not truth.

A gigantic biased dataset can produce an extremely precise estimate of the wrong quantity.

5. Uncertainty must be communicated, not hidden.

A bare estimate:

$$\hat{\theta} = 83.2$$

is often incomplete.

6. The unit of sampling matters.

If your 1,000 observations are actually 10 clusters of 100 highly correlated observations, your effective information may be much smaller than $n = 1000$ suggests.

7. Statistical inference is not a substitute for experimental design.

A sophisticated p -value cannot repair:

- bad sampling,
 - leakage,
 - confounding,
 - biased measurement,
 - invalid controls.
-

17. ONE-PAGE CONCEPT CONNECTION


TARGET POPULATION / PROCESS



SAMPLE DATA

Sample statistics	
mean	$\bar{x}$
SD	s
proportion	$\hat{p}$

SAMPLING DISTRIBUTION



STANDARD ERROR

Uncertainty interval
Confidence interval

HYPOTHESIS TESTING

SCIENTIFIC CLAIM

Interpretation + limits

18. WEEK 6 INTEGRATED WITH PREVIOUS WEEKS

This is where the course is becoming a coherent methodology.

Week 2

You learned:

Induction ↔ Deduction

and:

$$H \rightarrow P.$$

Week 4

You learned:

Question $\rightarrow$ Hypothesis $\rightarrow$ Prediction $\rightarrow$ Experiment $\rightarrow$ Measurement.

Week 5

You learned:

What makes hypotheses scientifically useful?

and:

How should measurement and reproducibility be approached?

Week 6

Now:

How uncertain is the measurement, and how do we infer population-level claims from finite data?

Thus:

Scientific reasoning + Measurement + Sampling + Probability = Statistical



This is a major conceptual milestone in the course.

19. KNOWLEDGE GAPS

Essential now

1. Confidence intervals

You should now understand:

- what they are,
 - what they are not,
 - how z and t intervals differ,
 - what confidence level means.
-

2. Sampling distributions

You should be able to explain:

Why does $\bar{X}$ have a distribution at all?

3. Standard error

You should be able to distinguish:

$$SD(X)$$

from:

$$SE(\bar{X}).$$

4. t -distribution

Understand:

- why it exists,
 - when it is used,
 - degrees of freedom,
 - why it converges to normal.
-

Useful later

5. Statistical power

You will need this for proper sample-size planning.

6. Type I and Type II errors

These naturally follow from the hypothesis-testing material.

7. Multiple comparisons

Especially important for AI experimentation.

Advanced / optional

8. Bootstrap inference

Very useful for AI metrics where analytical assumptions are inconvenient.

9. Bayesian credible intervals

Useful for understanding the contrast with frequentist confidence intervals.

10. Dependence / effective sample size

Extremely important for modern AI and time-series research.

20. FINAL REVISION SHEET

MUST KNOW

- Population vs sample
 - Parameter vs statistic
 - Sampling distribution
 - Central Limit Theorem
 - Standard error
 - Confidence interval
 - Error bar
 - z -score
 - Student's t
 - Degrees of freedom
 - Significance level
 - Bernoulli distribution
 - Binomial distribution
 - Sample proportion
 - Proportion standard error
-

MUST UNDERSTAND

1.

$$SE(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

means:

repeated sample means become more tightly concentrated as n increases.

2.

$$\hat{p} = \frac{X}{n}$$

is simply the sample mean of a 0/1 variable.

3.

$$SE(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}$$

is the sampling variability of the observed proportion.

4.

The CLT concerns the **sampling distribution**, not the population distribution.

5.

Confidence level describes the long-run performance of an interval procedure.

MUST REMEMBER

$$E[\bar{X}] = \mu$$

$$Var(\bar{X}) = \frac{\sigma^2}{n}$$

$$SE(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

$$t = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$$

$$df = n - 1$$

$$E[Y] = p$$

$$Var(Y) = p(1 - p)$$

$$X \sim \text{Binomial}(n, p)$$

$$E[X] = np$$

$$Var(X) = np(1 - p)$$

$$\hat{p} = \frac{X}{n}$$

$$E[\hat{p}] = p$$

$$Var(\hat{p}) = \frac{p(1 - p)}{n}$$

COMMON TRAPS

- ✗ "The CLT means the population becomes normal."
- ✗ " $n = 25$ always guarantees normality."
- ✗ "Standard deviation and standard error are the same."
- ✗ "A 95% CI means 95% probability that the fixed parameter is inside."
- ✗ "A p -value is the probability that H_0 is true."
- ✗ "Failing to reject H_0 proves H_0 ."
- ✗ "More data always eliminate uncertainty."
- ✗ "More data eliminate bias."
- ✗ "Every error bar means 95% confidence."
- ✗ "A binomial proportion can always be treated as normal."
- ✗ "1000 correlated observations equal 1000 independent observations."

10-MINUTE REVISION

If you revisit Week 6 months later, remember:

1.

We observe:

$$\bar{x}, s, \hat{p}$$

but care about:

$$\mu, \sigma, p.$$

2.

Different samples give different estimates.

Therefore the estimator has a:

sampling distribution.

3.

For the sample mean:

$$SE = \frac{\sigma}{\sqrt{n}}.$$

4.

For large enough n , the sampling distribution of the mean is approximately normal under suitable conditions.

5.

Confidence intervals turn sampling uncertainty into an interval estimate.

6.

For small samples with unknown σ , use the t -framework when its assumptions are appropriate.

7.

For binary data:

$$Y \in \{0, 1\}.$$

Then:

$$E[Y] = p,$$

$$\text{Var}(Y) = p(1 - p).$$

8.

For n binary observations:

$$X \sim \text{Binomial}(n, p).$$

9.

Observed proportion:

$$\hat{p} = X/n.$$

10.

Researcher question:

What assumptions make my uncertainty calculation trustworthy?

21. QUESTIONS WORTH INVESTIGATING FURTHER

These are the genuinely important next-level questions raised by this week:

1. How large does n actually need to be for the CLT approximation to be good?
2. What happens when observations are correlated?
3. How does statistical power determine sample size?
4. Why exactly does the t -distribution arise mathematically?
5. Why are confidence intervals not probability statements about fixed parameters?
6. When should we use a Wilson or exact binomial interval rather than the simple normal approximation?
7. How should confidence intervals be constructed for AI metrics such as F1, AUROC, BLEU, ROUGE or retrieval metrics?
8. What happens when the benchmark examples are dependent?
9. How should multiple hypothesis tests be corrected when running dozens of AI ablations?
10. How does Bayesian inference reinterpret uncertainty about μ or p ?

These are not merely side topics; several are likely to become important as the course progresses.

22. KNOWLEDGE BASE UPDATE

New Concepts

- Sampling distribution
- Central Limit Theorem
- Standard error
- Confidence interval
- Error bar
- z -score
- Student's t -distribution
- Degrees of freedom
- Significance level
- Bernoulli variable
- Binomial distribution
- Population proportion
- Sample proportion

- Proportion standard error

Reinforced Concepts

From earlier weeks:

- measurement,
- sampling,
- uncertainty,
- hypothesis testing,
- prediction,
- reproducibility,
- statistical reasoning,
- scientific inference.

Week 6 takes the measurement concepts of Week 5 and gives them a **formal probabilistic framework**.

Important Distinctions

Retain these long-term:

Population parameter $\neq$ Sample statistic

Population distribution $\neq$ Sampling distribution

$SD \neq SE$

Confidence level $\neq P(\text{parameter lies in this interval})$

p-value $\neq P(H_0 \mid \text{data})$

Fail to reject $\neq$ prove

More data $\neq$ less bias

25 $\neq$ universal CLT threshold

$\hat{p} \neq p$

Research Habits Acquired

1. Always identify the population being inferred about.
2. Identify the actual sampling unit.
3. Distinguish random variability from systematic problems.
4. Ask what the error bars actually represent.
5. State uncertainty alongside estimates.
6. Check assumptions before applying statistical formulas.
7. Never treat a statistical threshold as proof.

8. Avoid overclaiming from finite samples.

Future Dependencies

Likely important later:

- hypothesis testing,
 - statistical significance,
 - Type I/II errors,
 - power,
 - sample-size calculation,
 - regression,
 - uncertainty propagation,
 - experimental design,
 - possibly ANOVA/correlation.
-

23. CONNECTIONS TO YOUR AI RESEARCH JOURNEY

The biggest lesson for your AI work is:

An experimental result is not just a number; it is an estimate produced by a sampling and measurement process.

When you write:

"Our model achieves 91.4%."

you should mentally expand that to:

"Under this dataset, sampling procedure, evaluation protocol, model configuration and measurement definition, our estimate of performance is 91.4%, with some associated uncertainty."

That shift is extremely important for moving from:

AI implementation

toward:

AI research.

For your future paper reproductions, RAG experiments, SimCSE reproduction, and model comparisons, start thinking in terms of:

Claim → Target population → Sample → Estimator → Sampling variability



That is the statistical backbone of an empirical research paper.

24. RECOMMENDED STUDY ACTIONS

1. **Do the CLT numerically once.** Generate a strongly skewed population and repeatedly sample means for $n = 5, 10, 25, 50, 100$. Watch the sampling distribution evolve.
2. **Derive**

$$\text{Var}(\bar{X}) = \sigma^2/n$$

yourself. Don't merely memorise $SE = \sigma/\sqrt{n}$.

3. **Work through the $n = 9$, particle-mass t -test example by hand** until you can explain why the conclusion is "insufficient evidence," not "the mass is 21."
4. **Take one of your AI benchmark experiments and identify its population, sample, statistic, sampling unit and uncertainty source.**
5. **Next, learn statistical power and sample-size calculation.** That is the natural next conceptual step after understanding standard error and confidence intervals.

Sources



ChatGPT can make mistakes. Check important info.